A multivariate longitudinal model with subject-varying effects for the physical and mental progression of Alzheimer's Disease

Jesse W. Birchfield, Robert E. Weiss, and Andrew J. Holbrook

November 26, 2023

1   Introduction

Symptoms of Alzheimer's Disease (AD) include both physical and mental decline. Physically, by the prevailing theory [ref], amyloid plaques and neurofibrillary tangles progressively destroy the brain's neurons, beginning in the entorhinal cortex (ERC) and then spreading into the hippocampus and elsewhere [ref]. Mentally, subjects begin having difficulty with short-term memory and word recall, and then progressively lose other cognitive and motor abilities [ref]. In principle, the mind may be explicable in terms of the nervous system. In practice, it is not, and may never be. For the description of "objective" biological processes of the brain on one hand, and the description of "subjective" experiences and behaviors on the other, we have concepts, tools, and measurements that are not only different but incommensurable. A full account of the progression of AD, then, should include both kinds of description.

Whereas our previous work had the practical aim of improving cortical thickness estimates and enabling earlier detection, the present work is more purely exploratory. We investigate how the effects of the disease co-evolve over time, in relation to the covariates of age, sex, and diagnostic group (disease stage). Using a dataset from the Alzheimer's Diseases Neuroimaging Initiative (ADNI), in particular the ADNI-1 study [ref], we build a multivariate longitudinal model that includes both physical and mental outcomes for each subject at each timepoint. As a physical outcome, we use the thickness of the entorhinal cortex, as estimated by one or more imaging techniques. As a mental outcome, we use the measure taken by ADNI, performance on the Mini Mental State Examination (MMSE) [ref]. Our model features hierarchical structure, as varying slopes and intercepts for each subject and outcome (we prefer this term to "random effects"), with the means and full unstructured covariance matrix of these effects as hyperparameters. This construction enables maximal information sharing across outcomes and subjects, and it allows inference the mean and variance of any of the varying parameters, and the covariance between any subset of them. We discuss all of this in more depth below.

Immediately, we face an array of challenges. First is the challenge of imaging. For brain measurements, we start with magnetic resonance images (MRI), then use different computational "pipelines" to assemble them into a three-dimensional images, "segment" them into distinct tissue types, "label" them as different structures, then estimate their volume and thickness. Although sophisticated, these imaging pipelines are still subject to measurement error. Between pipelines, we see different estimates from the same scan. Within the same pipeline, between scans of the same individual, we see variation beyond that which can explained biologically. The magnitude of measurement error is large enough that it threatens to drown out the subtle effects we hope to detect statistically. And there is no ground truth or gold standard by which to measure the accuracy of these pipelines; we are limited to metrics of internal consistency and predictive power. See our previous work [ref] for an attempt to address this issue.

Second is the challenge of psychometrics. Next to the technological marvels of MRI machines and imaging algorithms, a brief questionnaire like the MMSE is very simple and cheap to administer and score. But it is difficult to determine what, exactly, is the construct being

measured, or what real entity a score (or score difference) corresponds to, or how accurate the instrument is as an estimator of that construct, or how accurate is an individual instance of measurement. The problem of no ground truth is more pronounced than for the task of imaging, because here it's hard to say even in theory what form ground truth would take. These are interesting and important issues, but beyond the scope of this research.

Third is the challenge of combining the two measures in a longitudinal model. Brain decay, for a single individual, does not occur uniformly throughout the organ, or at a steady tempo [ref]. Across individuals, the spatial and temporal decay patterns vary as well, possibly dependent on factors both measured and unmeasured. We should not expect a simple or consistent pattern of mental decline, either. Even the correlation between physical and mental symptoms is not straightforward. For example, subjects may silently experience several years of brain damage before manifesting any cognitive impairment [ref], and subjects with similar levels of brain damage may perform very differently on cognitive assessments [ref]. We propose a model flexible enough to capture this complexity and heterogeneity across outcomes, across time, and across subjects.

Our different outcomes also have very different numerical properties. Consistent with the hypothesis that they are generated as the true value plus a possible constant bias plus a random measurement error, ERC thickness estimates are positive real numbers with a "smooth", unimodal, symmetric distribution. A Gaussian fits easily. As questionnaire scores, however, MMSE values are integers, thus discrete and "coarse"; bounded, on a scale of 0 to 30; and with a very pronounced mode at the right boundary (30) and a long left tail. A Gaussian could hardly do worse. Still, the multivariate normal model has such clear advantages that we would like to adapt our data to the model, rather than the other way around. We propose a transformation of the MMSE data to make this possible.

We don't just specify a single model *a priori*. After some experimentation with model structure, we build eight models of the chosen class, with different configurations of ERC pipeline and MMSE transformation. We compare their predictive power with Bayesian 5-fold cross-validation. Since the models have different numbers of observations and outcomes on different scales, it takes some ingenuity to devise a metric comparable across all the models. Methodologically, this constitutes a novel combination of techniques and perhaps the most innovative aspect of this project.

[Interpretation: what do we learn from the model?]

2   Data

Our data comes from the Alzheimer's Disease Neuroimaging Initiative (Jack Jr et al., 2008; Petersen et al., 2010; Mueller et al., 2005), a public-private partnership sponsored in part by the National Institute on Aging (NIA). Beginning in October 2004, the three-year longitudinal ADNI-1 study followed a cohort of North American patients comprising 200 elderly controls, 400 with mild cognitive impairment (MCI), and 200 with diagnosed Alzheimer's Disease (AD). At each visit, researchers conducted cognitive assessments and structural MRI scans. Our analysis uses data from 663 of the 800 patients tabulated and made publicly avail-

able in Tustison et al. (2019), with a total of 2449 visits. We simply omit those visits that had missing MMSE scores.

[Description of FreeSurfer, FSLong, ANTs, SST. Why these two candidate pipelines?]

The MMSE is a screening test for cognitive impairment. It has 10 questions, takes 5 to 10 minutes to administer, and covers space and time orientation, naming familiar objects, repeating back a phrase, recalling a previous question, manipulating letters and/or numbers, and following instructions. Out of 30 possible points, a score of 24 or more indicates normal cognition, 19-23 indicates mild cognitive impairment, and 18 or below indicates moderate or severe impairment. While it can help differentiate different types of dementia, the test itself is not sufficient for diagnosis.

3   Methods

3.1   Rationale for model design

[rejected: ZINB model, interaction effects, hierachical modeling of fixed effect coefs] [explicate first model, then explain how the others differ]

3.2   MMSE transformation

The 2449 observations of MMSE look like this: [histogram of MMSE, normal qqplot] Clearly, they do not satisfy the normality assumption, and a multivariate normal model will give poor results.

A Gaussian distribution describes a continuous quantity, not a discrete one. MMSE score is literally a count variable–the number of questions answered correctly on the questionnaire. But since the MMSE's objective is to somehow quantify cognitive capacity, which can reasonably be thought of as continuous, we can treat MMSE score as a "coarse" or rounded measurement of a continuous latent variable.

Our first transformation is to consider "MMSE loss"–questions missed–rather than MMSE, by subtracting the raw value from 30. So a score of 30 becomes 0, a score of 29 becomes 1, and so on. We also imagine a kind of threshold, where cognitive decline has to reach a certain point before the subject would miss another question, so the MMSE loss represents a rounding *down* of the latent quantity. This justifies the step of adding 0.5 to each value. So define $y = 30 - MMSE + 0.5$. It has the following distribution in our data: [histogram of MMSE loss]

Our next step is much easier with a single-parameter probability distribution, and this looks like a histogram of a sample from some Exponential distribution. We assume $F(y) = 1 - e^{-\lambda y}$. The maximum likelihood estimate of λ is the maximum value of the polynomial $\Pi_{i=1}^{n} \left[e^{-\lambda y_i} - e^{-\lambda(y_i+1)} \right]$, which R easily computes as $\hat{\lambda} = 0.236$.

The probability integral transform states that if random variable W has continuous monotonic CDF F, then random variable $F(W)$ has a $Unif(0,1)$ distribution. Conversely, if

4

$U \sim Unif(0,1)$, and G is the continuous monotonic CDF of random variable X, then random variable $G^{-1}(U)$ has the same distribution as X. Letting F in the theorem be the exponential CDF with parameter $\hat{\lambda}$, and letting G be the standard normal CDF Φ, the transformation $\Phi^{-1}(F(Y))$ will have standard normal distribution, at least to the extent that F describes the distribution of Y. [histogram of transformed MMSE loss, hist of sample from expo(.236), QQplot] [histogram of MMSE loss]

It's true that this technique will give Y a *marginally* normal distribution, whereas the assumption of the multivariate normal regression model is that Y is *conditionally* normal given the predictors. But this is does not make much of a difference in practice. If $Y|X \sim N(X\beta, \sigma_Y^2)$, and $X \sim N(\mu_X, \sigma_X^2)$, then Y $N(\mu_X, \sigma_X^2 + \sigma_Y^2)$; it is still normal. If X is binary or discrete, then Y is a mixture of normals, empirically not much different from a true normal with the same mean and variance.

3.3   Model specification

Here we describe the model in prose, and afterward we formally specify it in symbols. For each subject, and at each timepoint, we have a bivariate or trivariate outcome, consisting of (i) ERC, as estimated by FSLong, ANTsSST, the arithmetic mean of the two, or both, and (ii) MMSE loss, either untransformed or transformed. We model these as bivariate or trivariate normal, where the mean is a linear function of predictors described shortly, and the error covariance matrix is unstructured, to be estimated from the data. Predictors are age at baseline, centered by subtracting the mean age at baseline across all 663 subjects, a gender indicator, and time of visit in years since baseline. For each outcome, we have a separate set of coefficients. The coefficients for age and gender are "fixed", in the sense that they do not vary by subject and do not have hyperparameters. The intercepts and slopes of the outcomes *do* vary by subject, and we consider them as draws from an underlying multivariate normal distribution with a common mean and an unstructured covariance matrix. This hierarchical structure allows for information sharing across outcomes, across timepoints, and across subjects. We do not add any additional covariance across timepoints other than that induced by the slopes. Finally, for each diagnostic group–CN, MCI, AD–we fit a completely separate set of coefficients and hyperparameters, so that we effectively have three separate regressions.

Formally, let i index subjects, $i = 1, ..., N$. For our data, $N = 663$. Let j index timepoints, $j = 1, ..., J_i$. For our data, J_i is between 2 and 6 for all i. Let k index outcomes, $k = 1, ..., K$. For models 1-3 and 4-6, $K = 2$. For models 4 and 8, $K = 3$. Let ℓ index groups, $\ell = 1, ..., L$. For our data, $\ell = 1, 2, 3$ for CN, MCI, AD respectively. Let m index covariates or coefficients. So we have five sets of indices i, j, k, ℓ, m, not all of which are relevant for any given quantity. To avoid confusion, we use dots as placeholders for irrelevant indices, so that $x_{7...2}$ is covariate $m = 2$ for subject $i = 7$, which is the same for all timepoints j and outcomes k and has group ℓ determined by i.

Let $\boldsymbol{y_{ij.\ell}}$ be the K-vector of observation ij, where subject i is in group ℓ. Let $\boldsymbol{\epsilon_{ij.\ell}}$ be the K-vector of errors for observation ij, for i in group ℓ. Both for computation and for inter-

pretability, we factor all covariance matrices into a diagonal matrix of standard deviations and a matrix of correlations, e.g. $C = SRS'$, and we write the standard deviations as a vector. Then, when specifying multivariate normal distributions, we use the notation $N(\text{mean}, \text{sd vector}, \text{correlation matrix})$. Let $\sigma_{...\ell}$ be the K-vector of standard deviations for the error vectors $\epsilon_{ij.\ell}$, and let $R_{...\ell}$ be the $K \times K$ correlation matrix.

Let $x_{i...1}$ be baseline age in years, mean-centered, for subject i. Let $x_{i...2}$ be the indicator that subject i is male, for subject i. Let $x'_{ij..} = \begin{bmatrix} x_{ij..1} & ... & x_{ij..M} \end{bmatrix}$. Generally, we use α to denote subject-invariant effects and β to represent subject-varying effects. Let $\alpha'_{..k\ell}$ be the 2-vector of coefficients for each $x'_{ij..}$ outcome k, group ℓ. Let $t_{ij...}$ be time in years since baseline for subject i, timepoint j. Let $z'_{ij...} = \begin{bmatrix} 1 & t_{ij..} \end{bmatrix}$. Let $\beta_{i.k\ell}$ be the 2-vector of subject-varying intercept and slope for subject i, outcome k, group ℓ. Let $\beta_{i..\ell}$ be the concatenated $2K$-vector $\begin{bmatrix} \beta'_{i.1\ell} & ... & \beta'_{i.K\ell} \end{bmatrix}'$. Let $\mu_{...\ell}$ be the $2K$-vector of means for the vectors $\beta_{i..\ell}$, let $\tau_{...\ell}$ be the $2K$-vector of standard deviations, and let $\Omega_{...\ell}$ be the $2K \times 2K$ correlation matrix.

This complexity of notation pays off, because it allows us to express the model quite compactly, as

$$
y_{ij.\ell} = \begin{bmatrix} x'_{ij..}\alpha_{..1\ell} + z'_{ij..}\beta_{i.1\ell} + \epsilon_{ij1\ell} \\ x'_{ij..}\alpha_{..2\ell} + z'_{ij..}\beta_{i.2\ell} + \epsilon_{ij2\ell} \\ x'_{ij..}\alpha_{..3\ell} + z'_{ij..}\beta_{i.3\ell} + \epsilon_{ij3\ell} \end{bmatrix}
$$

$$
\beta_{i..\ell} \mid \mu_{...\ell}, \tau_{...\ell}, \Omega_{...\ell} \sim N_4\left(\mu_{...\ell}, \tau_{...\ell}, \Omega_{...\ell}\right)
$$

$$
\epsilon_{ij.\ell} \mid \sigma_{...\ell}, R_{...\ell} \sim N_2\left(0, \sigma_{...\ell}, R_{...\ell}\right).
$$

We choose our priors to be as uninformative as possible while still allowing decent convergence in posterior sampling. We keep them identical across models and groups, and virtually identical across outcomes and coefficients. All prior distributions are independent of each other unless otherwise stated. Error sd's have prior $Expo(1)$, parameterized by rate. All α parameters have prior $N(0, 2)$. For ERC outcomes, means of subject-specific intercepts have prior $N(7, 3)$ and sd's have $Expo(1)$, while means of subject-specific slopes have prior $N(0, 1)$ and sd's have $Expo(5)$. For MMSE outcomes, means of subject-specific intercepts have prior $N(0, 2)$ and sd's have $Expo(1)$, while means of subject-specific slopes have prior $N(0, 2)$ and sd's have $Expo(1)$. Finally, all correlation matrices have prior $LKJCorr(1)$. This is the Lewandowski-Kurowicka-Joe distribution, a single-parameter family of correlation matrices (Gelman et al., 1995), argued to be superior to the traditional inverse Wishart prior distribution [ref: Stan manual]. A prior parameter $\eta = 1$ indicates neutrality about whether the correlation is positive, negative, or 0.

Our eight candidate models have the same structure, differing only in the outcome variables, as shown in Table 1:

Table 1: Candidate models

	y_1	y_2	y_3
1	FSLong	untransformed MMSE	
2	ANTsSST	untransformed MMSE	
3	ECTavg	untransformed MMSE	
4	FSLong	ANTsSST	untransformed MMSE
5	FSLong	transformed MMSE	
6	ANTsSST	transformed MMSE	
7	ECTavg	transformed MMSE	
8	FSLong	ANTsSST	transformed MMSE

3.4   Computation

3.5   Model comparison

4   Discussion

5   Appendix A: Stan code

```
// GENERAL SCRIPT FOR 8 MODELS

data{

  int<lower=1> I; // number of subjects
  int<lower=1> N; // total number of observations, each n is ij
  array[N] int<lower=1> id; // id of nth observation
  array[I] int<lower=1> group; // diagnostic group of subject i (CN=1, MCI=2
  int<lower=1> K; // number of outcomes
  array[N] vector[K] y; // vector of outcomes (e.g. FSLong, ANTsSST, MMSE_lo
  array[N] row_vector[4] X; // 1, t, age_c, male

}

transformed data {}

parameters{

  array[3] vector[2*K] alpha;
  array[3] vector[2*K] mu_beta;
  array[3] vector<lower=0>[2*K] sigma_beta;
  array[3] vector<lower=0>[K] sigma_epsilon;
  array[3] cholesky_factor_corr[2*K] L_Rho_beta;
  array[3] cholesky_factor_corr[K] L_Rho_epsilon;
```

```
    array[I] vector[2*K] z;

}

transformed parameters{

  array[N] vector[K] mu_y;
  array[I] vector[2*K] beta;
  array[3] cholesky_factor_cov[2*K] L_Sigma_beta;
  array[3] cholesky_factor_cov[K] L_Sigma_epsilon;


  for(l in 1:3){
    L_Sigma_epsilon[l] = diag_pre_multiply(sigma_epsilon[l], L_Rho_epsilon[
    L_Sigma_beta[l] = diag_pre_multiply(sigma_beta[l], L_Rho_beta[l]);
  }

  for(i in 1:I){
    // equiv to beta[i] ~ multi_normal(mu_beta[group[i]], Sigma_beta[group[
    beta[i] = mu_beta[group[i]] + L_Sigma_beta[group[i]] * z[i];
  }

  if(K == 2){
    for(n in 1:N){
      mu_y[n] = [ X[n] * append_row(beta[id[n]][1:2], alpha[group[id[n]]][1
                  X[n] * append_row(beta[id[n]][3:4], alpha[group[id[n]]][3
    }
  }

  if(K == 3){
    for(n in 1:N){
      mu_y[n] = [ X[n] * append_row(beta[id[n]][1:2], alpha[group[id[n]]][1
                  X[n] * append_row(beta[id[n]][3:4], alpha[group[id[n]]][3
                  X[n] * append_row(beta[id[n]][5:6], alpha[group[id[n]]][5
    }
  }

}

model{

// Likelihood

  for(n in 1:N){
    y[n] ~ multi_normal_cholesky(mu_y[n], L_Sigma_epsilon[group[id[n]]]);
```

```
  }

  // Priors
  for(l in 1:3){

    alpha[l] ~ normal(0,2);
    mu_beta[l][1] ~ normal(7,3); // y1 intercept mean
    mu_beta[l][2] ~ normal(0,1); // y1 slope mean
    sigma_beta[l][1] ~ exponential(1); // y1 intercept sd
    sigma_beta[l][2] ~ exponential(5); // y1 slope sd

    if(K == 2){
      mu_beta[l][3] ~ normal(0,2); // MMSE intercept mean
      mu_beta[l][4] ~ normal(0,2); // MMSE slope mean
      sigma_beta[l][3] ~ exponential(1); // MMSE intercept sd
      sigma_beta[l][4] ~ exponential(1); // MMSE slope sd
    }

    if(K == 3){
      mu_beta[l][3] ~ normal(7,3); // y2 intercept mean
      mu_beta[l][4] ~ normal(0,1); // y2 slope mean
      mu_beta[l][5] ~ normal(0,2); // MMSE intercept mean
      mu_beta[l][6] ~ normal(0,2); // MMSE slope mean
      sigma_beta[l][3] ~ exponential(1); // y2 intercept sd
      sigma_beta[l][4] ~ exponential(5); // y2 slope sd
      sigma_beta[l][5] ~ exponential(1); // MMSE intercept sd
      sigma_beta[l][6] ~ exponential(1); // MMSE slope sd
    }

    L_Rho_beta[l] ~ lkj_corr_cholesky(1);
    sigma_epsilon[l] ~ exponential(1);
    L_Rho_epsilon[l] ~ lkj_corr_cholesky(1);

  }

  // random effects distribution
  for(i in 1:I){
    z[i] ~ std_normal();
  }

}

generated quantities {

  vector[N] log_lik;
```

```stan
  array[3] matrix[2*K,2*K] Rho_beta;
  array[3] matrix[K,K] Rho_epsilon;

  for(n in 1:N){
    log_lik[n] = multi_normal_cholesky_lpdf(y[n] | mu_y[n], L_Sigma_epsilon
  }

  for(l in 1:3){
    Rho_beta[l] = L_Rho_beta[l] * L_Rho_beta[l]';
    Rho_epsilon[l] = L_Rho_epsilon[l] * L_Rho_epsilon[l]';
  }

}
```